

Neural Networks: Architecture, Types, and Training

Report of research on Neural Networks (600 – 750 words)

The landscape of technology has been transformed by neural networks, which are one of the key foundations for contemporary AI as it is applied to tasks that encompass everything from language translation to autonomous vehicle navigation. Inspired by the structure of the human brain, these artificial complex systems are composed of interconnected nodes (neurons) that work in parallel to process data and learn patterns. To tap into the potential of neural networks in solving complex problems, you need to understand a lot about them, such as basic architecture, types of networks, and how training happens.

At the core of any neural network are its neurons arranged in a number and sequence of (not necessarily unique) layers. Neural networks consist of three layers: the input layer, hidden layers (if any), and output layer. The input layer is the main entry or loading area for data in our neural system, and each neuron serves as a container that can represent any distinguishable attribute of the distinctive feature from incoming information subsets. Intermediate layers (referred to variably as hidden) of weights are subsequently applied, with their number and size determined by hyperparameters, which derive from the data through backpropagation during training. Lastly, the output layer is responsible for generating what will be its final prediction or classification as performed by a network.

Every neuron in all these layers is doing a simple but especially important job — or computational step. It accepts some input data, processes it through a mathematical function known as an activation function, and returns the output. This processing is governed by the activation function, a key element that adds non-linearity to the network so it can learn complex patterns in data. Popular activation functions include the sigmoid function, which maps all input values to and (ReLU) that just outputs zero for negative numbers. An example of this is the (ReLU) function, which has become standard due to its ability to address the vanishing gradient problem that often occurs during deep neural network training.

Neural networks are not one piece of monolithic brick; you have various kinds for relevant tasks. Three types: Feedforward neural networks (FNNs) are the simplest type, in which data flows from your input layer to any hidden layers and then into an output layer only. This simplicity makes FNNs great for basic classification and regression. But for more complex work, those general architectures will not cut it.

One example of this is convolutional neural networks (CNNs), which are tailored to image processing tasks. CNNs are excellent at capturing spatial hierarchies with convolutional layers that multiply input data by a filter — identifying edges, textures, and similarly useful artifacts for image recognition or object detection. CNNs have revolutionized the field of computer vision, delivering astounding successes in areas ranging from medical image analysis to autonomous vehicles.

Recurrent Neural Networks (RNNs) are another type of neural network that is especially fit for sequential data. This is an advantage of RNNs over feedforward networks, which have connections going in only one direction and thus are unable to remember anything about their past. This unique property makes RNNs well-suited for tasks involving sequences like time-series prediction, NLP (Natural Language Processing), etc. However, RNNs notoriously face the issue of learning long-term dependencies, and more advanced architectures like Long Short-Term Memory (LSTM) networks or Gated Recurrent Units (GRUs) have been designed to improve this aspect.

A neural network is trained by iteratively adjusting the weights of its connections between units to cancel out that difference. This is mostly done through an algorithm called backpropagation, where the gradient of the loss function (error) against each weight in a network is calculated. The first step is in the training process, where your model goes through bytes of input data to produce an output. The resulting output is compared to the actual target values, and a loss would be calculated.

The backward pass: Backpropagation is the entire process of computing gradients $dL/d(\text{weight})$ for each weight using the chain rule, to be able to perform stochastic gradient descent. These gradients tell us the direction and size by which each weight should change to decrease this error. After the gradients are calculated, optimization algorithms are used to update weights. Stochastic Gradient Descent (SGD) is one of the most widely used optimizers where weights are updated in small batches, instead of all the data at once. Adam (Adaptive Moment Estimation) and its many variants also improve SGD by adapting learning rates on a per-parameter basis, allowing more efficient training speed.

Regularization is a common way to improve training with the help of optimization techniques, but there are many ways of doing this and another popular method used along with improvements using regularization. Regularization, such as L2 regularization and dropout, helps prevent the model from overfitting when it becomes too closely fitted to the data, resulting in deficient performance on unseen examples.

Over the last few years, neural networks have proven to be an incredibly powerful tool in a wide variety of problems, from computer vision through NLP. This same ability to learn complexity and generalize from data is what has made them so essential in modern AI. A deeper understanding of the architecture, types, and training process helps identify what neural networks can do for us in AI research, thus helping push forward the boundaries further. There is no question; we are going to see neural nets play a much more significant role in the future of tech as R&D continues. Mongo DB vs SQL.

References

<https://inria.hal.science/hal-01637477v1/document>

<https://www.ncbi.nlm.nih.gov/books/NBK583971/>

<https://www.v7labs.com/blog/neural-networks-activation-functions>

<https://www.ibm.com/topics/neural-networks>

<https://aws.amazon.com/what-is/neural-network/>

<https://encord.com/blog/activation-functions-neural-networks/>